

Exercises 02582
Module 1
Spring 2018

January 31, 2018

Topics: OLS - Ridge - Bias - Variance

Resources for this exercise:

Listing 1: Resources in Matlab

```
readtable('DiabetesDataNormalized.txt') % load data in the .txt file
      % (from CampusNet)
y = T(:,11); % response
X = T(:,1:10); % matrix with variables
whos % lists variables in memory – check what you have loaded (X and y)
rand(n, p); % simulate random uniform data
randn(n, 1); % simulate random gaussian data
logspace(-4, 3, 100) % 100 values on the logarithmic scale
      % from 1e-4 to 1e3
knnsearch(X, X(i,:), 'K', K+1) % find the K nearest neighbours.
```

Listing 2: Resources in R

```
library('lars') # load lars library (includes diabetes data)
data(diabetes) # load data
str(diabetes$x) # list structure of x
str(diabetes$y) # list structure of y
runif(n*p, min = 0, max = 1) # simulate random uniform data
rnorm(n, mean=0, sd = 1) # simulate random normal data
library('pracma') # load pracma library for 'logspace'
logspace(-4, 3, 100) # 100 values on the logarithmic scale from 1e-4 to 1e3
get.knnx(X, query, k=K+1, ...) # find the K nearest neighbors
```

Listing 3: Resources in Python

```
import numpy as np # numpy library
import scipy.linalg as lng # linear algebra from scipy library
import matplotlib.pyplot as plt # library for plots
from scipy.spatial import distance # load distance function
from sklearn import preprocessing as preproc # load preproc function

T = np.loadtxt('DiabetesDataNormalized.txt', delimiter = '_',
skiprows = 1) # load data in the .txt file (from CampusNet)
y = T[:, 10] # response
X = T[:, :10] # variables
np.random.randn(n, p) # simulate random gaussian data
np.random.rand(n, 1) # simulate random uniform data
np.logspace(-4, 3, k) # 100 values on the logarithmic scale
# from 1e-4 to 1e3
preproc.scale(X) # normalize to zero mean and unit variance
distance.euclidean(X[i, :], X[j, :]) # compute distance between two points
```

Exercises 02582
Module 1
Spring 2018

January 31, 2018

Topics: OLS - Ridge - Bias - Variance

Exercises:

- 1 Solve the Ordinary Least Squares (OLS) computationally (for the diabetes data set):
 - (a) What is the difference between using a brute force implementation for an OLS solver and a numerically 'smarter' implementation? Compute the ordinary least squares solution to the diabetes data set for both options and look at the relative difference.
 - (i) In Matlab: Use for example `inv` and Matlab's backslash operator (for info type `help mldivide`).
 - (ii) In R: Use for example `solve` to invert the matrix or to solve the linear system of equation.
 - (iii) In Python: Use for example `linalg.lstsq` to invert the matrix or to solve the linear system of equation.
 - (b) How could you include an intercept term in Matlab/R/Python? This means using the model $\mathbf{y} = \beta_0 + \mathbf{x}_1\beta_1 + \dots + \mathbf{x}_p\beta_p + \epsilon$ rather than $\mathbf{y} = \mathbf{x}_1\beta_1 + \dots + \mathbf{x}_p\beta_p + \epsilon$.
 - (c) Calculate the mean squared error $MSE = 1/n \sum_{i=1}^n (y_i - \mathbf{x}_i\boldsymbol{\beta})^2$.
 - (d) Calculate the residual sum of squares $RSS = \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|_2^2$ and the total sum of squares $TSS = \|\mathbf{y} - \bar{y}\|_2^2$, where \bar{y} is the estimated mean of \mathbf{y} . Report on the R^2 measure, that is, the proportion of variance in the sample set explained by the model: $R^2 = 1 - \frac{RSS}{TSS}$.

2 Examine Bias and Variance for the OLS:

- (a) Investigate the unbiasedness of OLS using simulated data as follows.
 - (i) Create a random matrix \mathbf{X} consisting of three random variables which are NID, with $X \sim N(\mathbf{0}, \mathbf{I})$ and sample 10 observations from each.
 - (ii) Create the true regression parameters $\boldsymbol{\beta}_{true} = [1, 2, 3]^T$.
 - (iii) Create the response \mathbf{y} by the linear model $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$. Make sure the errors $\boldsymbol{\epsilon}$ are homoscedastic and zero-mean $\epsilon \sim N(0, \sigma^2)$, where σ^2 denotes the noise level in the simulated data.
 - (iv) Estimate the regression coefficients $\boldsymbol{\beta}$ from this data
 - (v) Repeat steps (iii)-(iv) 100 times.
 - (vi) Use meaningful plots to investigate bias and variance for the model. Experiment with different noise levels σ^2 .

3 Solve the Ridge regression problem and examine Bias and Variance for Ridge:

- a) Derive (using pen and paper) the ridge regression solution by, as you would when minimizing any differentiable analytical function, differentiating $\|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\| + \|\boldsymbol{\beta}\|_2$ with respect to $\boldsymbol{\beta}$, setting to zero and solving for $\boldsymbol{\beta}$. That is, solve $\frac{\partial}{\partial \boldsymbol{\beta}} [\|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|_2^2 + \lambda \|\boldsymbol{\beta}\|_2^2] = 0$ for $\boldsymbol{\beta}$.
- b) Compute ridge-regression solutions for the diabetes data set for 100 values of the regularization parameter in the range 10^{-4} to 10^3 . Plot the solutions as a function of this parameter. In the next lecture you will learn how to choose a single parameter value which suits the problem at hand.
- c) Change the experiment in exercise 2 to investigate bias and variance for ridge regression instead of OLS. Can you lower the variance without introducing too much bias?

4 Implement and solve KNN regression:

- (a) Implement a KNN regression algorithm:
 - (i) Find the K nearest neighbours using a suitable distance metric (e.g. Euclidean).
 - (ii) Optional: Compute weights for the neighbours as the proportion of its distance to the total distance for the K nearest neighbours.
 - (iii) Compute the predicted response as the (weighted) mean of the K neighbours.
- (b) Find a solution to the diabetes data using KNN regression. Try different options for K .

End of exercise